

Seminar

July 25th, 2020

DATA  
SCIENCE  
INDONESIA

Data Science  
Indonesia  
East Java Region

# JATIM CAMP #5

Build Data Ecosystem for Better Analytics

## Data Pipelining for Creating Great Data Ecosystem with



**Rendy Bambang Junior**  
Senior Data Manager, Ruangguru



**Muhammad Iqbal Tawakal**  
Data Scientist, Gojek

Supported by



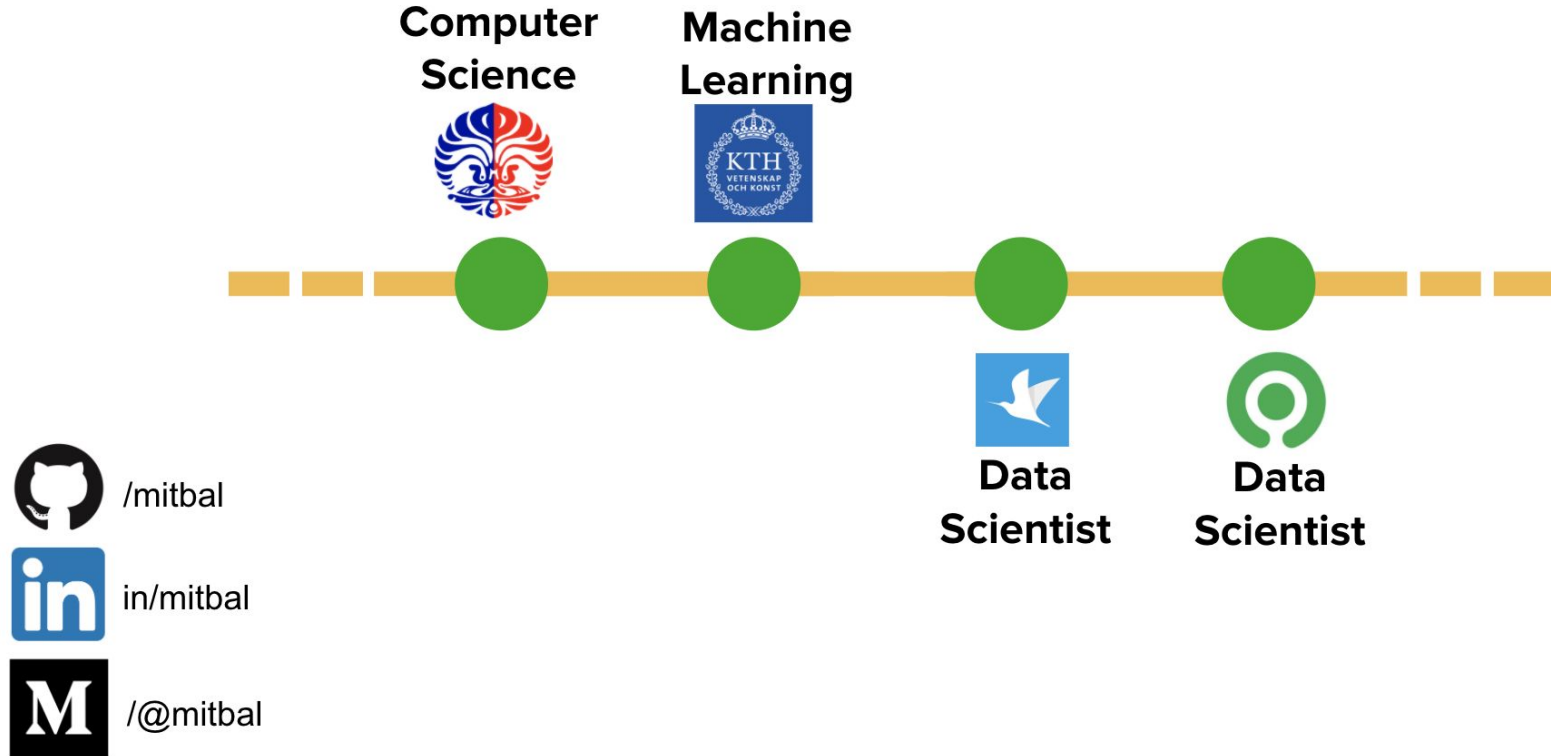


# Practical Data Pipeline for Data Science

Case study: BigQuery in Gojek

M Iqbal Tawakal (@mitbal)  
DS@gojek

# Self (Re)Introduction



# Outline

- Introduction to Gojek, and its DS team
- Intro to DS workflow and challenges faced
- How proper data pipeline can help alleviate them
- How BigQuery fit into all this
- Pros, cons, and other consideration

# Gojek in Southeast Asia

Operates in **207 cities** in  
**Southeast Asia**



**>155  
millions**  
App  
Downloads



**+400  
thousands**  
Merchant  
Partners (96%  
SMEs)



**+2 millions**  
Driver  
Partners

# Range of products and **solutions**

 goride

 gocar

 gosend

 gobox

 gobluebird

 gofood

 gobuy

 goshop

 gomed

 gotix

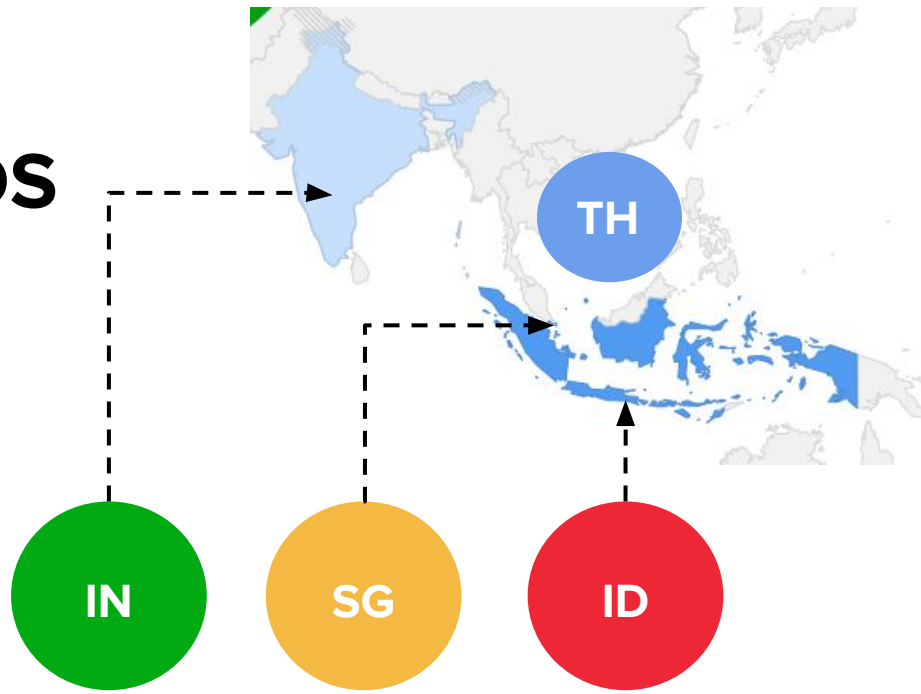
 gopay

 gopoints

 gopulsa

 gobills

# Geo @ DS

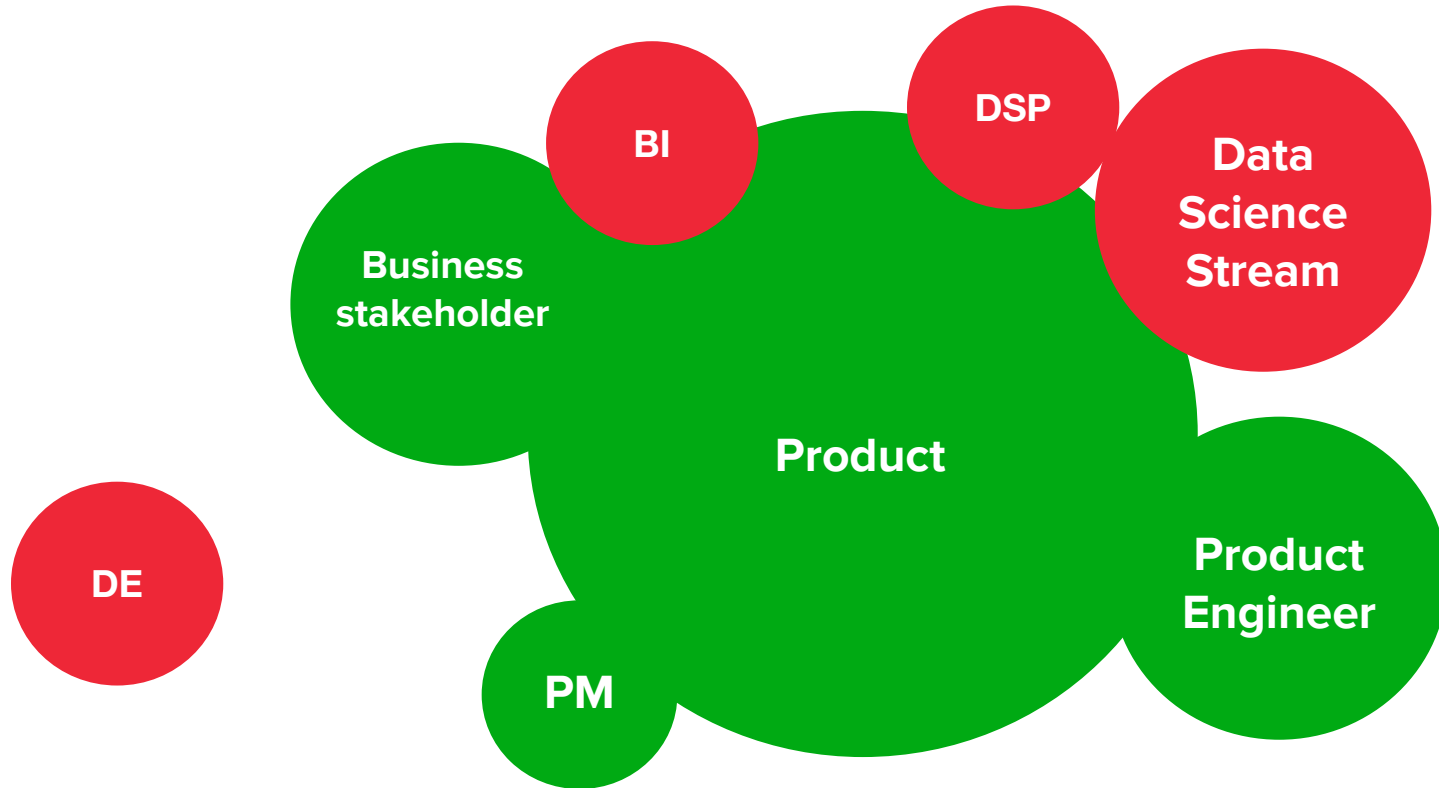


# Streams @ DS

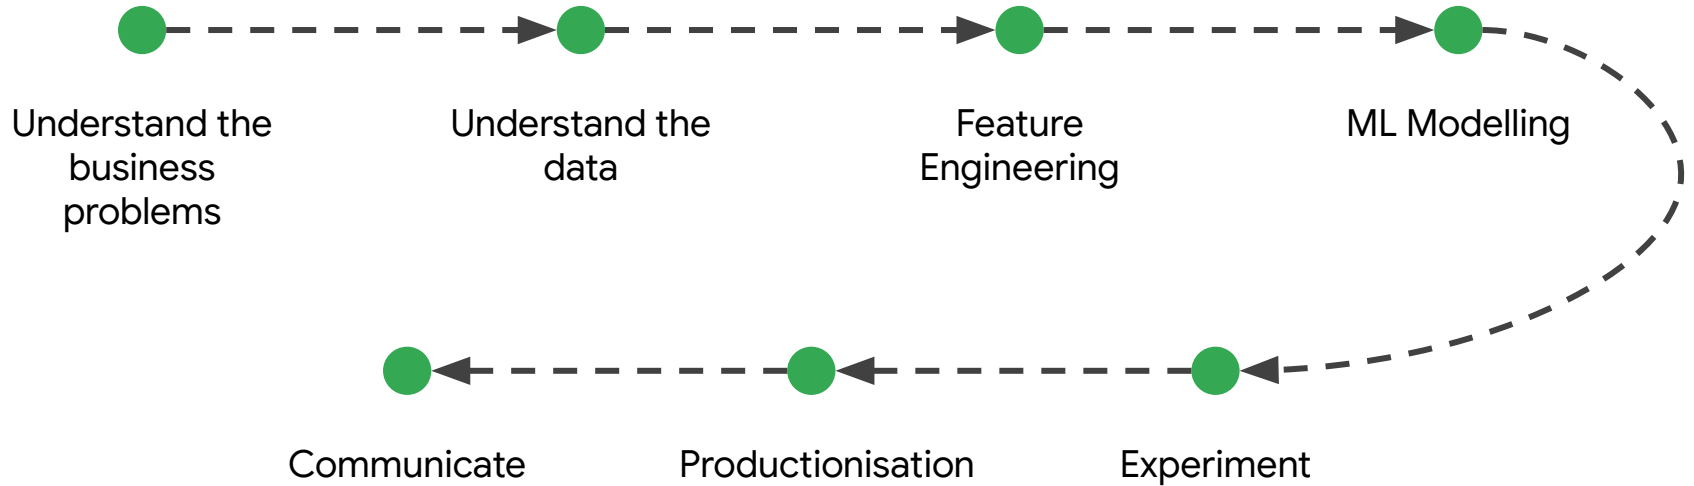




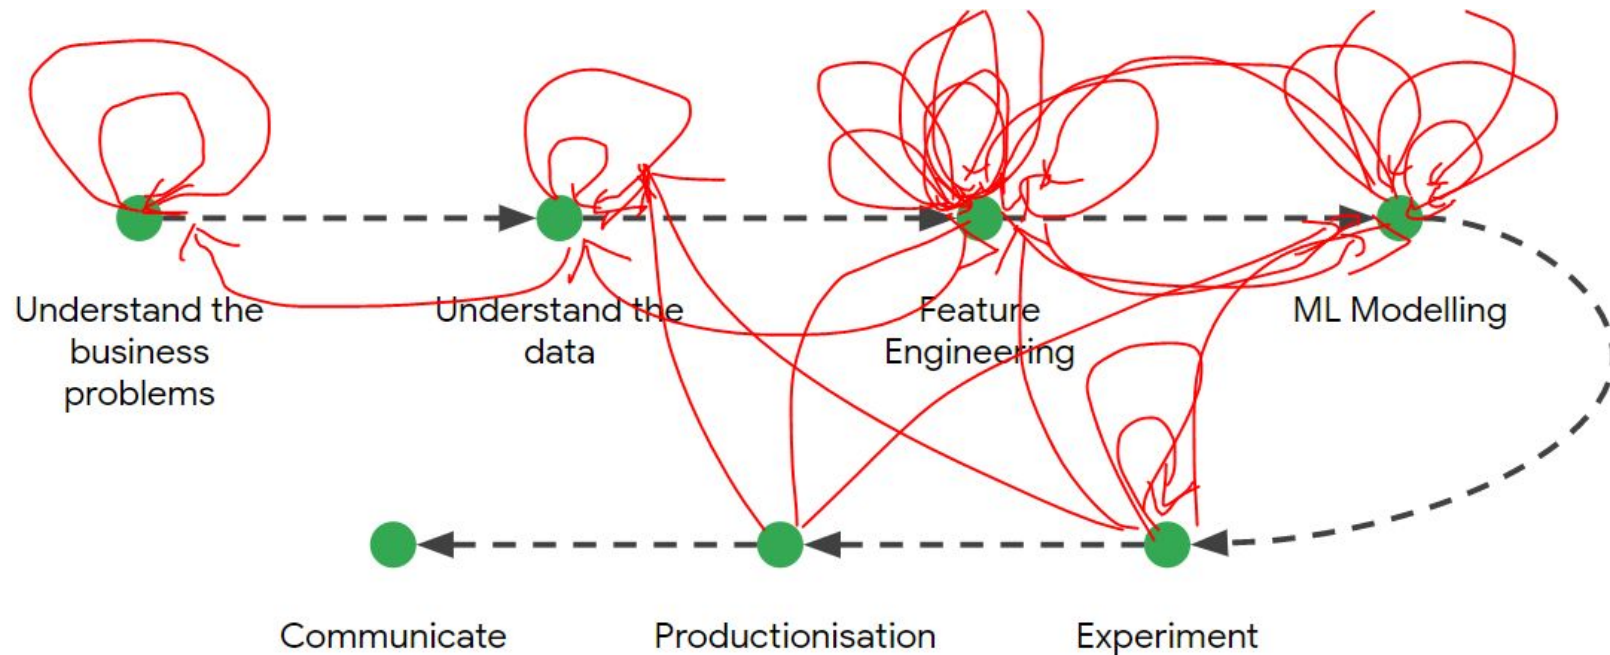
# DS, Embedded to Product



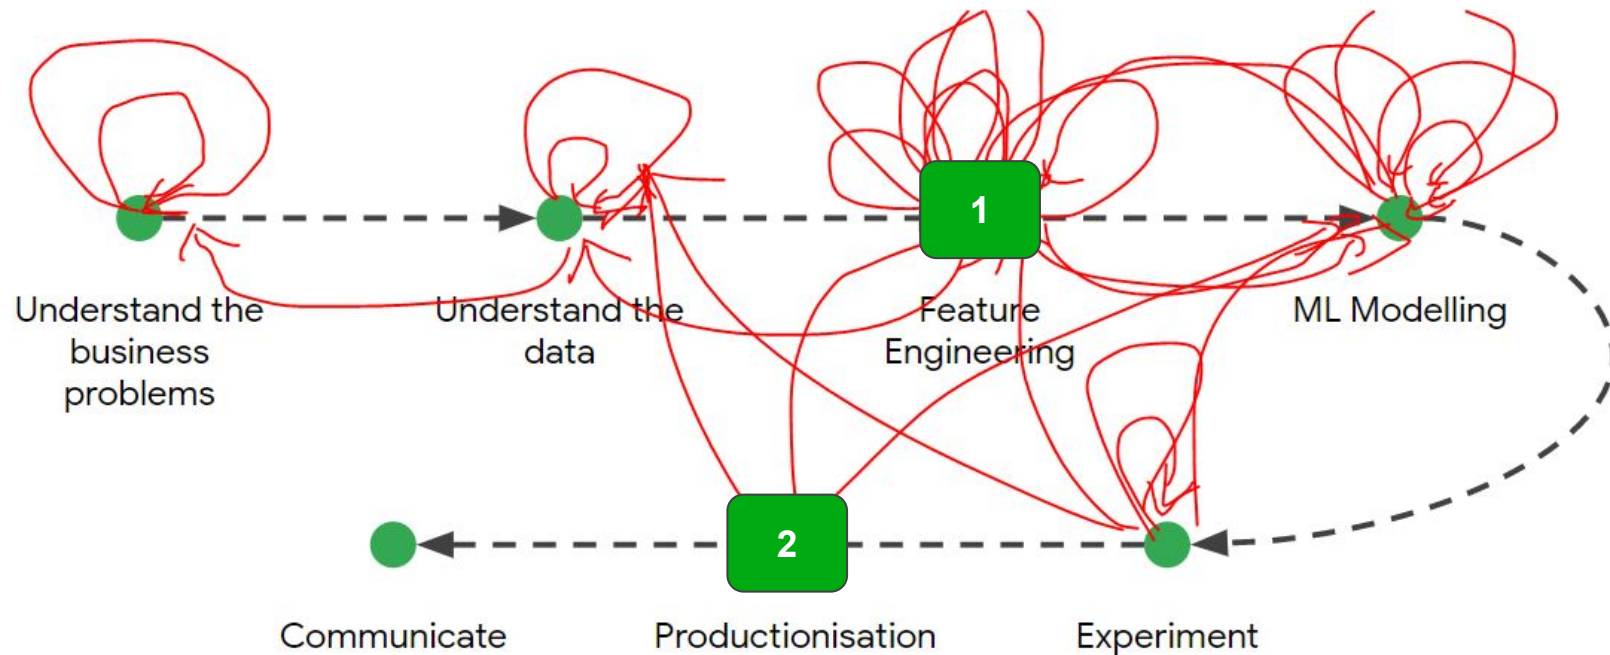
# (Ideal) Data Science Timeline and Workflows



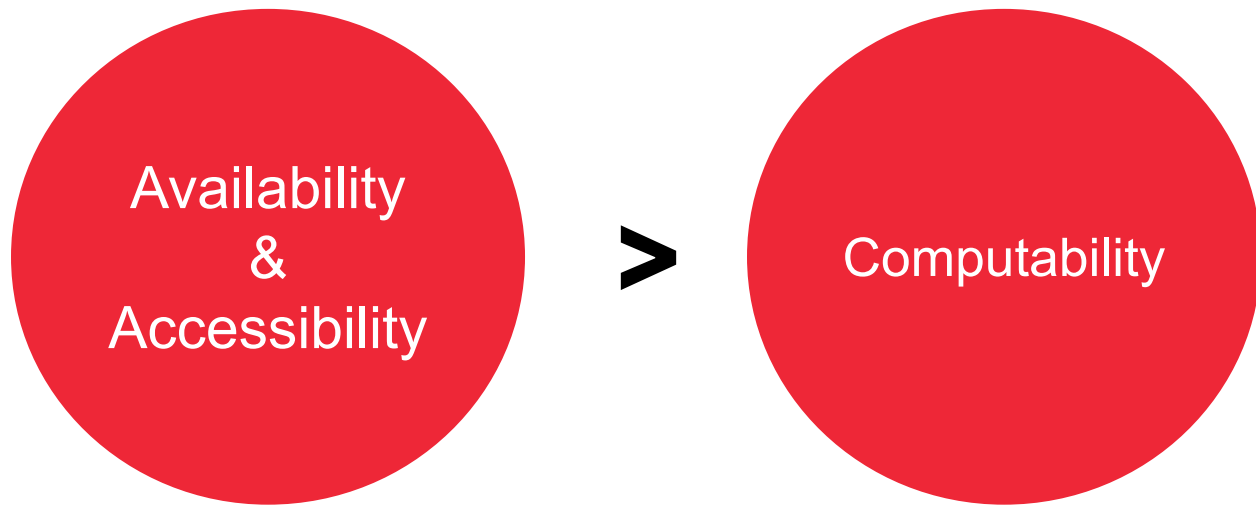
# Actual Data Science Workflow!!?!?!



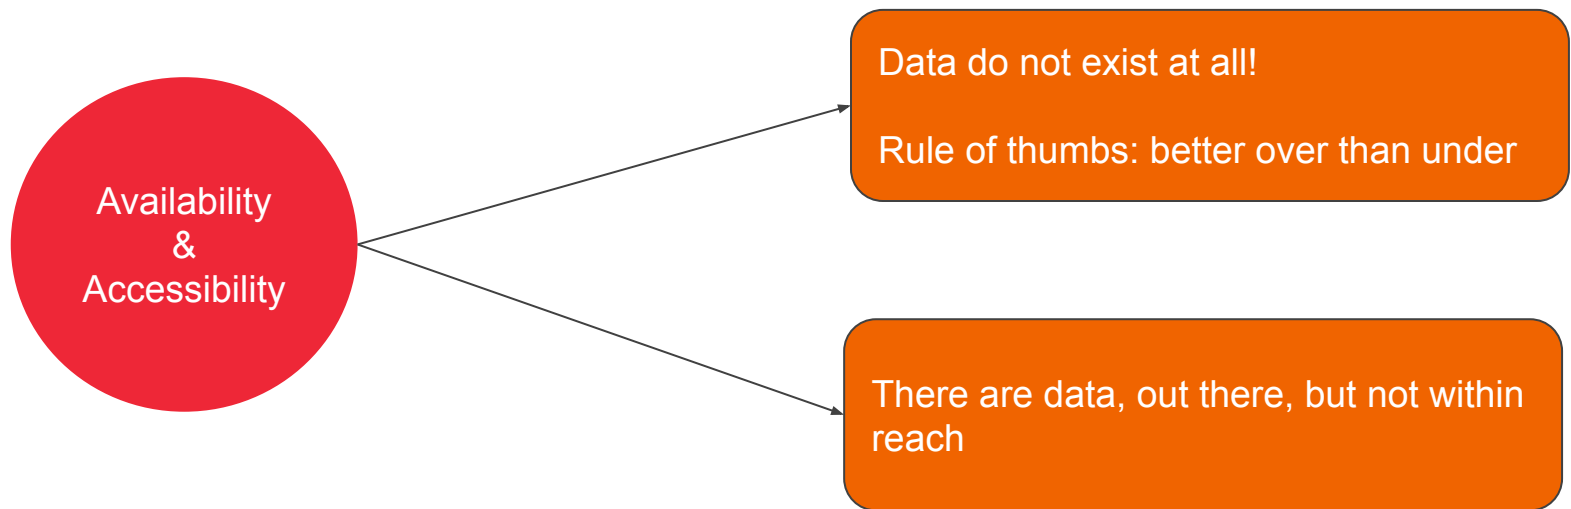
# Where proper DS pipeline can help most



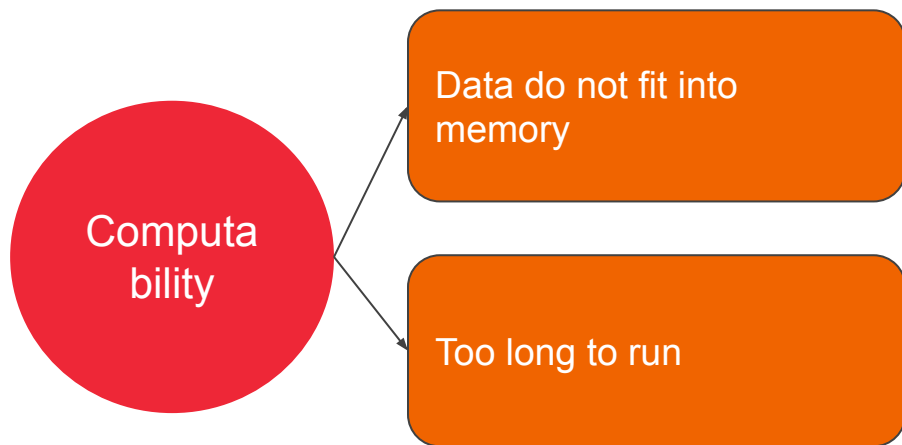
# 1st part: 2 core challenges



# Availability & Accessibility Problem



# Computability Problem



# BigQuery come into Picture

## BigQuery

Serverless, highly scalable, and cost-effective multi-cloud data warehouse designed for business agility.



# Not the only player in town



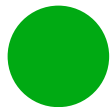
Processing engine alternative



Data warehouse alternative



# For your consideration



## Pros

- combine storage & processing engine in one package
- no need to manage physical or virtual instance
- computation can scale with data size automatically
- using SQL as the interface language to interact

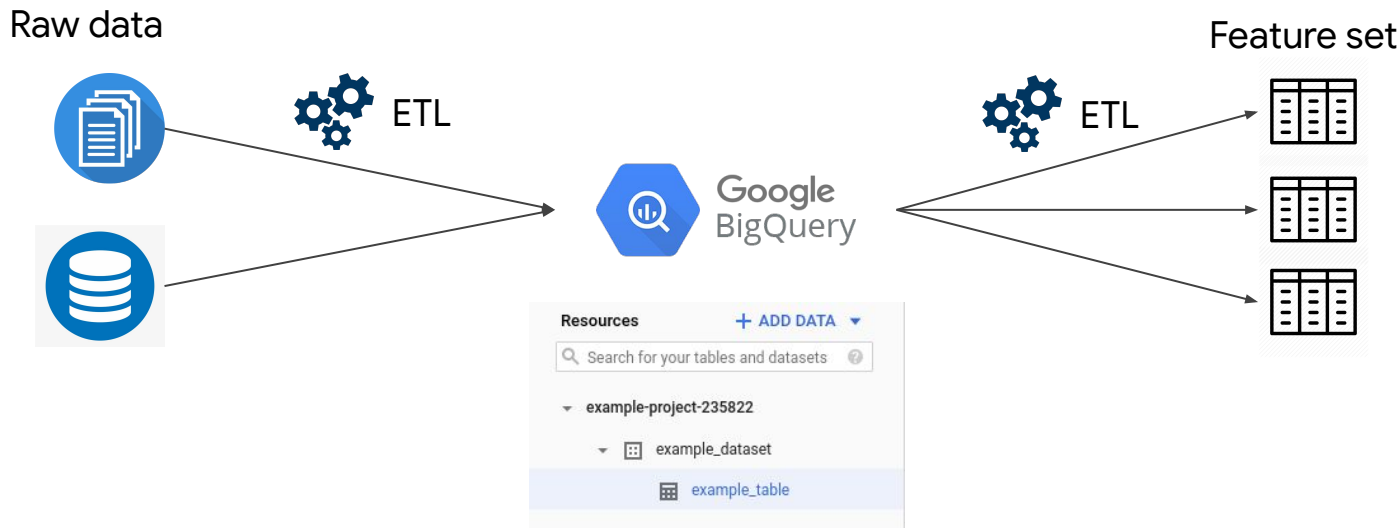


## Cons

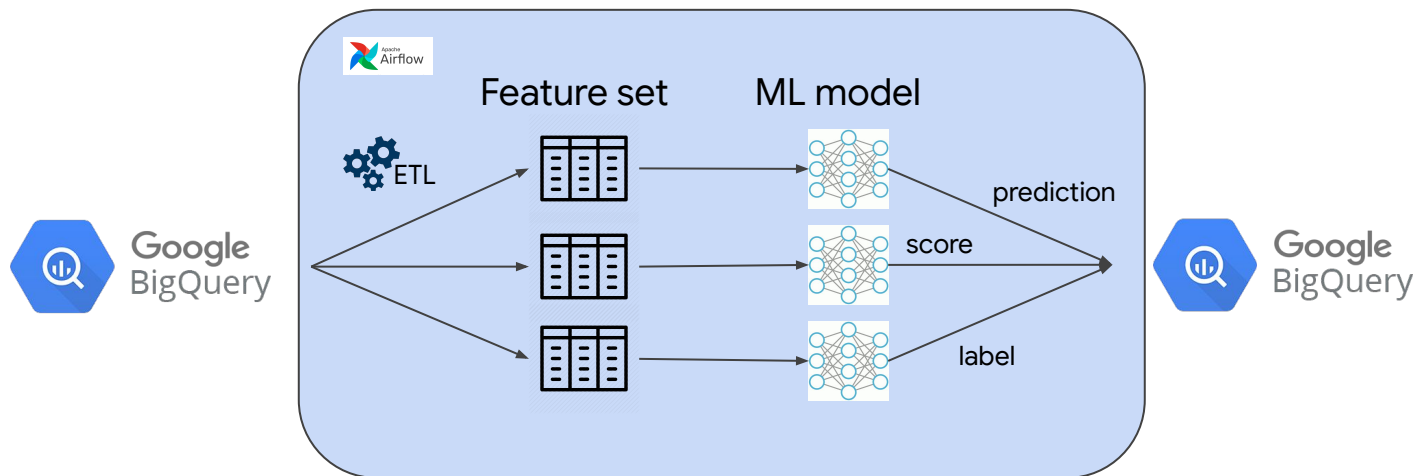
- can be costly if usage is not monitored properly
- might not suited for certain types of operation

Operation	Pricing	Details
Active storage	\$0.020 per GB	The first 10 GB is free each month. See <a href="#">Storage pricing</a> for details.
Long-term storage	\$0.010 per GB	The first 10 GB is free each month. See <a href="#">Storage pricing</a> for details.
BigQuery Storage API	\$1.10 per TB	The BigQuery Storage API is not included in the <a href="#">Google Cloud Free Tier</a> .
Streaming Inserts	\$0.010 per 200 MB	You are charged for rows that are successfully inserted. Individual rows are calculated using a 1 KB minimum size. See <a href="#">Streaming pricing</a> for details.
Queries (on-demand)	\$5.00 per TB	The first 1 TB per month is free. See <a href="#">On-demand pricing</a> for details.

# Accessibility solved?: 1 place to put them all



# Batch deployment option



# Computability solved?: Compute at any scale, up and down



Small OK

```
1 SELECT
2   ...,
3   COUNT(...)
4 FROM
5   ...
6 where
7   _partitiontime = '2020-06-22'
8 group by 1
9 order by 2 desc
```

Processing location: US

[Run](#) [Save query](#) [Save view](#) [Schedule query](#)

Query results

[SAVE RESULTS](#)

[EXPLORE DATA](#)

Query complete (1.2 sec elapsed, 7.1 MB processed)



Big No problem

```
1 SELECT
2   ...,
3   ...,
4   ...,
5   ...
6 FROM (
7   SELECT
8     ...,
9     EXTRACT(minute
10    FROM
11      datetime(event_timestamp,
12        'Asia/Jakarta')) AS minute,
```

Processing location: US

[Run](#) [Save query](#) [Save view](#) [Schedule query](#) [More](#)

Query results

[SAVE RESULTS](#)

[EXPLORE DATA](#)

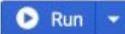
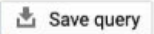
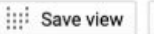


Query complete (4.0 sec elapsed, 120.2 GB processed)


# Best Practice: Do and Don't (1)

- Never select \*, explicitly choose field and attribute to be included into processing

 Before



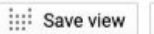


```
1 SELECT
2 *
3 FROM
4 [REDACTED]
5
```


    

This query will process 1.2 TB when run. 

 After

```
1 SELECT
2 [REDACTED],
3 [REDACTED]
4 FROM
5 [REDACTED]
6
```

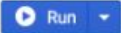


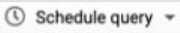

This query will process 18.1 GB when run. 


# Best Practice: Do and Don't (2)

- Always filter by partition date

 Before

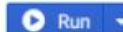




```
1 SELECT
2 ...
3 ...
4 FROM
5 ...
```


    

This query will process ~~18.1 GB~~ when run. 

 After

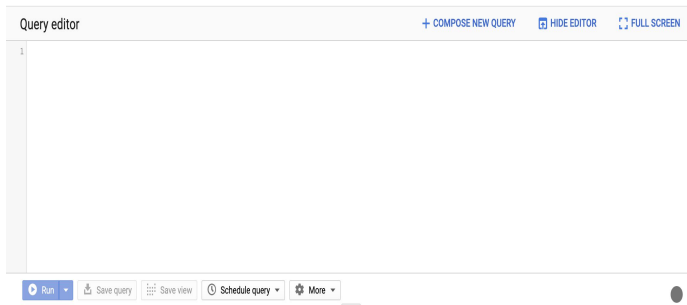
```
1 SELECT
2 ...,
3 ...,
4 FROM
5 ...
6 WHERE
7   _partitiontime = '2020-02-22'
```

This query will process ~~17.3 MB~~ when run. 

# How to use as data scientist

## Web UI



- Autoformat
- Syntax highlighting
- Save result

## Programmatically

```
In [ ]: !pip install pandas_gbq
```

executed in 3.36s, finished 00:42:43 2020-07-22

```
In [ ]: query = """
        select field1, field2
        from `project_name.dataset_name.table_name`
        where _partitiontime = '2020-07-25'
        """
```

```
In [ ]: df = pd.read_gbq(query, dialect='standard')
```

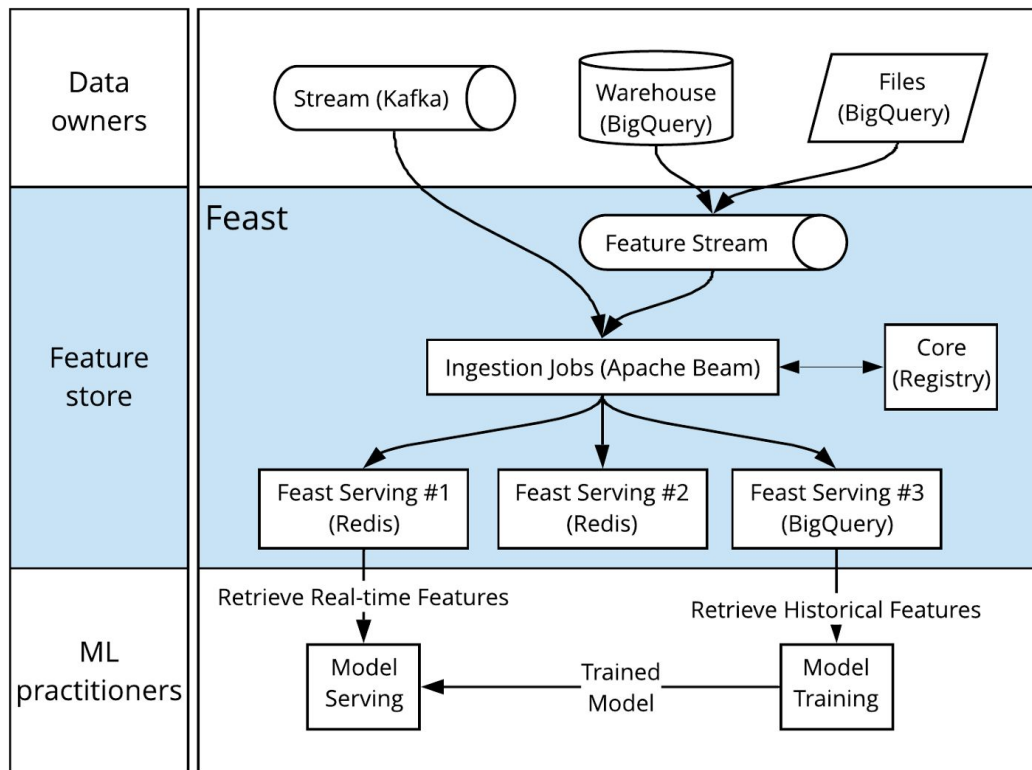
- Need to setup authentication first
- Seamlessly used in the next processing steps



# Going Beyond: Feature engineering reusability and serving with Feature Storage

Platform	Open-Source	Offline	Online	Metadata	Feature Engineering	Supported Platforms	TimeTravel / Point-in-Time Queries	Training Data
<a href="#">Hopsworks</a>	AGPL-V3	Hudi/Hive	MySQL Cluster	DB Tables, Elasticsearch	(Py)Spark, Python	AWS, GCP, On-Prem	SQL Join or Hudi Queries	.tfrecords, .csv, .npy, .petastorm, .hf5, etc
<a href="#">Michelangelo</a>	N/A	Hive	Cassandra	KV Entries	Spark, DSL	Proprietary	SQL Join	Streamed to models?
<a href="#">Feast</a>	Apache V2	BigQuery	BigTable/Redis	DB Tables	Beam, Python	GCP	SQL Join	Streamed to models
<a href="#">Conde Nast</a>	N/A	Kafka/Cassandra	Kafka/Cassandra	Protocol Buffers	Shared libraries	Proprietary	?	Protobuf
<a href="#">Zipline</a>	N/A	Hive	KV Store	KV Entries	Flink, Spark, DSL	Proprietary	Schema	Streamed to models?

# FEAST: Gojek own Feature Storage solution



# Closing Remark

---

1

Data Science and Machine Learning project workflow rarely work in straightforward manner. Good data pipeline is almost become necessity for it to be successful

---

2

BigQuery can be used to fill this shoes. However, with big power comes big responsibility. Is it the only final solution? No, but it's a good start

---

3

Need strong data foundation pipeline ready first

---

4

Data science is a team sport. All of these cannot happen without, the biz with their infinite wisdom, DSP with their cool tools, product engineering for their system, and BI for maintaining original data pipeline

**Thank You!**

